

Assisty Grading Assistant

The Assisty Grading Assistant scores student submissions — case analyses, short-answer and essay exams, applied assignments — against your marking scheme, and returns a defensible mark, a supporting quote, and personalised feedback for every student. It can also grade the open-ended work that answer-key tools cannot, and it flags its own uncertain grades for you to review. You stay in control of the rubric and the final grade and feedback.

What makes it different

Personalised feedback at scale. Each student receives per-question strengths and specific areas to improve, in plain language you can forward with minimal editing — the same quality for the 200th student as the 1st.

Handles various submission formats. The tool can grade submissions in a variety of formats: pdf, doc, docx, ppt, pptx, html, ipynb notebooks, plain-text and also scanned hand-written submissions. Handwritten submissions too poor to read reliably are flagged for manual checking rather than silently mis-graded.

Can grade open-ended work. It marks against behavioural rubric levels, not a single 'correct' solution, so it works for case studies, essays, and applied analysis where good answers vary.

Rubrics are generated. From a one-line marking scheme it drafts a detailed five-level rubric for you to review — tuned against a sample of how your students actually answered.

One rubric, applied identically to everyone. The same standard is applied to every script across every section and TA, removing grader-to-grader drift and 'who marked my paper' disputes.

It flags when it is unsure. A grading that the grader is unsure of is flagged as low-confidence and routed to you. You review the doubtful few, not all of them.

Marks are defensible. Each score cites the exact sentence from the student's own work that justifies it. If there is no supporting evidence, the criterion scores zero — the tool never awards unearned marks.

You control the strictness. Choose a lenient, moderate, or strict marking style and see the whole cohort re-scored, so you set the standard deliberately rather than discovering it after release.

How faculty use it — user stories

Grade the pile

US-1 As a professor with 200 student assignment submissions, I want to upload a batch and receive scored, evidenced results in hours, so that I return grades in days, not weeks. *[bulk ingest + batch grading]*

US-2 As faculty grading pen-and-paper exams, I want scripts the tool couldn't read clearly flagged rather than mis-graded, so that I only hand-check the unreadable few. *[OCR quality triage]*

Consistency & fairness

US-3 As a coordinator running four sections with four TAs, I want every script held to one identical standard, so that a grade doesn't depend on who marked the paper. *[single calibrated rubric]*

US-4 As a professor facing a grade appeal, I want every mark backed by a quote from the student's own answer, so that I can resolve the appeal in minutes with evidence. *[evidence-cited grading]*

US-5 As a program director, I want to re-score a batch at lenient, moderate, and strict settings, so that I can set the curve on purpose. *[adjustable strictness]*

Feedback & teaching insight

US-6 As an instructor, I want every student to get specific, per-question feedback, so that the whole cohort gets guidance I could never write by hand at scale. *[auto feedback]*

US-7 As a professor, I want to see which questions the cohort scored lowest on, so that I can focus on teaching that concept in the next class. *[cohort analytics]*

Trust & adoption

US-8 As a cautious first-time user, I want the tool to auto-grade the confident cases and surface only the uncertain ones, so that I spend my time where it disagrees with itself. *[confidence review queue]*

US-9 As a department head evaluating the tool, I want to compare its grades against marks my TAs already gave, so that I can measure agreement before trusting it live. *[human-in-the-loop validation]*

Frequently Asked Questions

What faculty most often ask before adopting the Grading Assistant

Q. Is this going to replace my judgment as an examiner?

A. No. It is an assistant, not an autopilot. You define the marking scheme, review and edit the rubric it drafts, and you hold the final grade. The intended workflow is that it does the first pass on the whole batch and flags the cases it is unsure about, so your expertise is spent on the hard calls and final sign-off rather than on repetitive marking.

Q. How accurate is it, and how do I know I can trust a grade?

A. Every script is graded three times independently. Where all three runs agree, the grade is marked high-confidence; where they disagree, it is flagged low-confidence and sent to you. Rather than asking you to trust a single opaque number, the tool tells you which grades are solid and which need a human eye — and backs each one with a direct quote from the student's work. We recommend a first run where you compare its grades against your own on a sample, so you can calibrate your trust before going live.

Q. Does it work for subjective, open-ended answers where there is no single right answer?

A. Yes — absolutely. It does not need to compare answers to a fixed solution key. Instead it grades against five behavioural performance levels per criterion (for example, from ‘fully addresses the point with justification’ down to ‘not addressed’), which is how a human examiner reads a case study or essay. This is why it suits business case analysis, strategy memos, and applied reasoning, not just questions with one fixed answer.

Q. Can I upload supplementary materials (like case studies or articles) for the assistant to reference?

A. Yes, absolutely. The setup process encourages you to prepare and upload all relevant assignment data along with supporting materials, such as specific articles or case studies, beforehand. This ensures the grading assistant has the necessary context to evaluate student answers accurately against the background material.

Q. Do I have to manually type out my questions and criteria to set up a new assignment?

A. Not necessarily. When creating a new assignment, you have two options:

- **Option 1 (Prefill):** If you already have an existing question file, you can upload it into the "Assignment source". Assisty will use this file to **automatically extract** the assignment details, questions, and marking criteria. *(Note: You should carefully review this auto-generated content before finalizing, as minor corrections may be needed.)*
- **Option 2 (Manual):** You can choose to enter your assignment details, questions, and marking scheme manually

Q. Do I have to write a full rubric myself?

A. No. You provide the question and a one-line marking scheme (‘what a full-credit answer contains’), and the tool expands it into a complete five-level rubric. It even calibrates the partial-credit levels against a sample of your students’ actual answers, so the middle bands reflect the real mistakes your cohort makes. You review and edit the draft before grading — the rubric is a transparent document, not a black box.

Q. Can I control how strict or lenient the marking is?

A. Yes. You choose a lenient, moderate, or strict marking style, and the whole cohort is scored accordingly. Because all three are computed together, you can compare the resulting grade distributions and decide the standard deliberately before releasing anything.

Q. What happens with messy handwriting or a poor-quality scan?

A. Each scanned script is automatically checked for readability. Anything the system cannot read with confidence is flagged as 'needs review' and set aside for you to grade manually, rather than being given an unreliable machine grade. It fails safe: an unreadable script is surfaced, never quietly guessed.

Q. Could the grader 'make things up' or award marks the student didn't earn?

A. The design guards against this directly. A criterion can only score above zero if the tool can quote the exact supporting text from the submission; if it finds no evidence, the criterion is forced to zero. Because every mark is tied to a real quote from the student's own answer, you can spot-check any grade against the source in seconds.

Q. What about students who used AI to write the answer, or plagiarism?

A. The tool grades the quality of the answer against your rubric; it is not a plagiarism or AI-detection service and does not claim to be. It fits alongside your existing integrity checks rather than replacing them. Its evidence quotes can, however, make it easier to notice answers that are off-topic or suspiciously generic.

Q. What subjects and assignment types does it support?

A. It is subject-agnostic. It has been used on business subjects such as Organizational Behaviour case exams, and it applies equally to technical and analytical courses. It suits any written, open-ended assessment: short-answer and essay exams, case analyses, applied problem write-ups, and report-style submissions, whether typed or handwritten.

Q. What kind of cohort performance analytics does the platform provide?

A. Once grading is complete, the analytics dashboard provides a detailed visual overview of performance, displaying:

- The highest, lowest, and average scores.
- The overall score distribution across the cohort.
- Lists of the top five and bottom five performers.

- Question-level analytics that break down and compare student performance on each individual item, making it easy to identify which concepts need re-teaching

Q. How much of my time does it actually save, and what is the setup effort?

A. The one-time effort is preparing the questions and marking schemes and reviewing the drafted rubric — work comparable to preparing a marking guide you would create anyway. After that, grading a full batch and generating per-student feedback runs unattended, and your remaining effort is concentrated on the flagged low-confidence scripts and final sign-off.

Q. If I disagree with a recommended grade, what specific adjustments can I make?

A. You remain the examiner of record and hold complete control over the final outcome. When reviewing an individual student's evaluation, you can select the edit option to directly:

- Modify the grading assistant's written feedback.
- Adjust the specific rubric criteria comments.
- Manually increase or decrease the assigned points for individual criteria.

Additionally, you can adjust the overall grading strictness on the fly by selecting from three levels: Lenient, Moderate, or Strict. Switching between these levels will automatically recalculate the student's total score based on the chosen standard.

Nothing is final until you approve it.

Q. Will students receive feedback, and can I edit it first?

A. Yes. For each student the tool drafts per-question strengths, specific areas to improve, and a short overall summary, written in plain language without internal grading jargon. It is yours to review and edit before it reaches the student.

Q. Is student data kept secure and private?

A. Student submissions are processed only to produce the grades and feedback you request, within your institution's approved infrastructure, and are not used to train public models. Exact data-handling, residency, and retention arrangements are confirmed as part of your institution's deployment — your IT and data-protection teams can review these before rollout.